

Multiple linear regression model for Bike sharing

July 30, 2022

1 PROBLEM GOAL:

Develop a model to find the variables that are significant in the demand for shared bikes with the available independent variables and report appropriate metrics of your model evaluation.

2 Data link:

<https://docs.google.com/spreadsheets/d/1WieH-IPBXkKPj46XCEsjAgShaRdZ3A9BTN2zz27jSN4/edit?usp=sha>

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

3 Reading data:

```
[2]: bike = pd.DataFrame(pd.read_csv("day.csv"))
```

```
[3]: bike.head()
```

```
[3]:
```

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	\
0	1	01-01-2018	1	0	1	0	1	1	
1	2	02-01-2018	1	0	1	0	2	1	
2	3	03-01-2018	1	0	1	0	3	1	
3	4	04-01-2018	1	0	1	0	4	1	
4	5	05-01-2018	1	0	1	0	5	1	

	weathersit	temp	atemp	hum	windspeed	casual	registered	\
0	2	14.110847	18.18125	80.5833	10.749882	331	654	
1	2	14.902598	17.68695	69.6087	16.652113	131	670	
2	1	8.050924	9.47025	43.7273	16.636703	120	1229	
3	1	8.200000	10.60610	59.0435	10.739832	108	1454	
4	1	9.305237	11.46350	43.6957	12.522300	82	1518	

cnt

```

0    985
1    801
2   1349
3   1562
4   1600

```

4 Check the information of the dataset:

```
[4]: bike.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 730 entries, 0 to 729
Data columns (total 16 columns):
 #   Column        Non-Null Count  Dtype  
---  -
 0   instant      730 non-null    int64  
 1   dteday       730 non-null    object  
 2   season       730 non-null    int64  
 3   yr           730 non-null    int64  
 4   mnth         730 non-null    int64  
 5   holiday      730 non-null    int64  
 6   weekday      730 non-null    int64  
 7   workingday   730 non-null    int64  
 8   weathersit    730 non-null    int64  
 9   temp         730 non-null    float64 
10  atemp        730 non-null    float64 
11  hum          730 non-null    float64 
12  windspeed    730 non-null    float64 
13  casual       730 non-null    int64  
14  registered   730 non-null    int64  
15  cnt          730 non-null    int64  
dtypes: float64(4), int64(11), object(1)
memory usage: 91.4+ KB

```

5 To check the descriptive information of the dataset:

```
[5]: bike.describe()
```

```

[5]:
count    instant    season    yr    mnth    holiday    weekday  \
count    730.000000    730.000000    730.000000    730.000000    730.000000    730.000000
mean     365.500000     2.498630     0.500000     6.526027     0.028767     2.995890
std      210.877136     1.110184     0.500343     3.450215     0.167266     2.000339
min        1.000000     1.000000     0.000000     1.000000     0.000000     0.000000
25%      183.250000     2.000000     0.000000     4.000000     0.000000     1.000000
50%      365.500000     3.000000     0.500000     7.000000     0.000000     3.000000
75%      547.750000     3.000000     1.000000    10.000000     0.000000     5.000000

```

max	730.000000	4.000000	1.000000	12.000000	1.000000	6.000000
-----	------------	----------	----------	-----------	----------	----------

	workingday	weathersit	temp	atemp	hum	windspeed \
count	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000
mean	0.690411	1.394521	20.319259	23.726322	62.765175	12.763620
std	0.462641	0.544807	7.506729	8.150308	14.237589	5.195841
min	0.000000	1.000000	2.424346	3.953480	0.000000	1.500244
25%	0.000000	1.000000	13.811885	16.889713	52.000000	9.041650
50%	1.000000	1.000000	20.465826	24.368225	62.625000	12.125325
75%	1.000000	2.000000	26.880615	30.445775	72.989575	15.625589
max	1.000000	3.000000	35.328347	42.044800	97.250000	34.000021

	casual	registered	cnt
count	730.000000	730.000000	730.000000
mean	849.249315	3658.757534	4508.006849
std	686.479875	1559.758728	1936.011647
min	2.000000	20.000000	22.000000
25%	316.250000	2502.250000	3169.750000
50%	717.000000	3664.500000	4548.500000
75%	1096.500000	4783.250000	5966.000000
max	3410.000000	6946.000000	8714.000000

6 To check the shape of the data:

```
[6]: print(bike.shape)
```

```
(730, 16)
```

```
[7]: bike.columns
```

```
[7]: Index(['instant', 'dteday', 'season', 'yr', 'mnth', 'holiday', 'weekday',
          'workingday', 'weathersit', 'temp', 'atemp', 'hum', 'windspeed',
          'casual', 'registered', 'cnt'],
          dtype='object')
```

7 Drop columns that are not useful for analysis:

Based on the high level look at the data and the data dictionary, the following variables can be removed from further analysis:

instant : Its only an index value

dteday : This has the date, Since we already have seperate columns for 'year' & 'month',hence, we could live without this column.

casual & registered : Both these columns contains the count of bike booked by different categories of customers. Since our objective is to find the total count of bikes and not by specific category,

we will ignore these two columns. More over, we have created a new variable to have the ratio of these customer types.

We will save the new dataframe as bike_new, so that the original dataset is preserved for any future analysis/validation

```
[8]: bike_new=bike[['season', 'yr', 'mnth', 'holiday', 'weekday',  
                  'workingday', 'weathersit', 'temp', 'atemp', 'hum', 'windspeed',  
                  'cnt']]
```

```
[9]: bike_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 730 entries, 0 to 729  
Data columns (total 12 columns):  
#   Column          Non-Null Count  Dtype  
---  ---  
0   season          730 non-null   int64  
1   yr              730 non-null   int64  
2   mnth            730 non-null   int64  
3   holiday         730 non-null   int64  
4   weekday         730 non-null   int64  
5   workingday      730 non-null   int64  
6   weathersit       730 non-null   int64  
7   temp            730 non-null   float64  
8   atemp           730 non-null   float64  
9   hum             730 non-null   float64  
10  windspeed       730 non-null   float64  
11  cnt             730 non-null   int64  
dtypes: float64(4), int64(8)  
memory usage: 68.6 KB
```

```
[10]: bike_new['season']=bike_new['season'].astype('category')  
bike_new['weathersit']=bike_new['weathersit'].astype('category')  
bike_new['mnth']=bike_new['mnth'].astype('category')  
bike_new['weekday']=bike_new['weekday'].astype('category')
```

```
C:\Users\MAHA\AppData\Local\Temp\ipykernel_3252\468937576.py:1:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
bike_new['season']=bike_new['season'].astype('category')  
C:\Users\MAHA\AppData\Local\Temp\ipykernel_3252\468937576.py:2:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.
```

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
bike_new['weathersit']=bike_new['weathersit'].astype('category')
C:\Users\MAHA\AppData\Local\Temp\ipykernel_3252\468937576.py:3:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
bike_new['mnth']=bike_new['mnth'].astype('category')
C:\Users\MAHA\AppData\Local\Temp\ipykernel_3252\468937576.py:4:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
bike_new['weekday']=bike_new['weekday'].astype('category')
```

```
[11]: bike_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 730 entries, 0 to 729
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   season          730 non-null   category
1   yr              730 non-null   int64
2   mnth            730 non-null   category
3   holiday         730 non-null   int64
4   weekday         730 non-null   category
5   workingday      730 non-null   int64
6   weathersit       730 non-null   category
7   temp            730 non-null   float64
8   atemp           730 non-null   float64
9   hum             730 non-null   float64
10  windspeed       730 non-null   float64
11  cnt             730 non-null   int64
dtypes: category(4), float64(4), int64(4)
memory usage: 49.7 KB
```

8 SPLITTING THE DATA

Splitting the data to Train and Test: - We will now split the data into TRAIN and TEST (70:30 ratio)

We will use `train_test_split` method from `sklearn` package for this

```
[12]: from sklearn.model_selection import train_test_split
      np.random.seed(0)
      df_train, df_test = train_test_split(bike_new, train_size = 0.70, test_size = 0.
      ↪30, random_state = 333)
```

```
[13]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 510 entries, 483 to 366
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   season      510 non-null   category
 1   yr          510 non-null   int64
 2   mnth       510 non-null   category
 3   holiday     510 non-null   int64
 4   weekday     510 non-null   category
 5   workingday  510 non-null   int64
 6   weathersit   510 non-null   category
 7   temp        510 non-null   float64
 8   atemp       510 non-null   float64
 9   hum         510 non-null   float64
10  windspeed   510 non-null   float64
11  cnt         510 non-null   int64
dtypes: category(4), float64(4), int64(4)
memory usage: 38.9 KB
```

```
[14]: df_train.shape
```

```
[14]: (510, 12)
```

```
[15]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 219 entries, 22 to 313
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   season      219 non-null   category
 1   yr          219 non-null   int64
 2   mnth       219 non-null   category
 3   holiday     219 non-null   int64
```

```

4  weekday      219 non-null    category
5  workingday   219 non-null    int64
6  weathersit    219 non-null    category
7  temp         219 non-null    float64
8  atemp        219 non-null    float64
9  hum          219 non-null    float64
10 windspeed    219 non-null    float64
11 cnt          219 non-null    int64
dtypes: category(4), float64(4), int64(4)
memory usage: 17.3 KB

```

```
[16]: #check the shape before splitting
df_test.shape
```

```
[16]: (219, 12)
```

```
[17]: #check the train set columns
df_train.columns
```


```
[17]: Index(['season', 'yr', 'mnth', 'holiday', 'weekday', 'workingday',
        'weathersit', 'temp', 'atemp', 'hum', 'windspeed', 'cnt'],
        dtype='object')
```

9 Outlier checking

```
[18]: #Draw boxplot for indepent variables
cols = ['temp', 'atemp', 'hum', 'windspeed']
plt.figure(figsize=(18,4))


i = 1
for col in cols:
    plt.subplot(1,4,i)
    sns.boxplot(y=col, data=bike_new)
    i+=1

```



10 Draw pair Plots to check the linear relationship

```
[19]: #Draw pairplots for continuous numeric variables using seaborn
bike_num=df_train[['temp', 'atemp', 'hum', 'windspeed', 'cnt']]
sns.pairplot(bike_num, diag_kind='kde')
plt.show()
```



11 Find the Correlation between the Numerical Variable

```
[20]: bike_new.corr()
```




```
[20]:
```

	yr	holiday	workingday	temp	atemp	hum	\
yr	1.000000	0.008195	-0.011852	0.048789	0.047215	-0.112547	
holiday	0.008195	1.000000	-0.257009	-0.028764	-0.032703	-0.015662	
workingday	-0.011852	-0.257009	1.000000	0.002044	0.010657	0.053770	
temp	0.048789	-0.028764	0.002044	1.000000	0.991696	0.128565	
atemp	0.047215	-0.032703	0.010657	0.991696	1.000000	0.141512	
hum	-0.112547	-0.015662	0.053770	0.128565	0.141512	1.000000	
windspeed	-0.011624	0.006257	-0.002453	-0.158186	-0.183876	-0.248506	
cnt	0.569728	-0.068764	-0.027640	0.627044	0.630685	-0.098543	


	windspeed	cnt
yr	-0.011624	0.569728
holiday	0.006257	-0.068764
workingday	-0.002453	-0.027640
temp	-0.158186	0.627044
atemp	-0.183876	0.630685
hum	-0.248506	-0.098543
windspeed	1.000000	-0.235132
cnt	-0.235132	1.000000

```
[21]: #Draw Heatmap
sns.heatmap(bike_new[['temp','atemp','hum','windspeed','cnt']].corr(),
            cmap='Blues', annot = True)
plt.show()
```




12 Analysing Categorical Variabels with target variables


```
[22]: plt.figure(figsize=(3,3),dpi=100)
plt.title("Cnt vs Temp",fontsize=10)
sns.regplot(data=bike_new,y="cnt",x="temp")
plt.xlabel("Temperature")
plt.show()
```



```
[23]: plt.figure(figsize=(3,3),dpi=100)
plt.title("Cnt vs Windspeed",fontsize=10)
sns.regplot(data=bike_new,y="cnt",x="windspeed")
plt.show()
```



```
[24]: num_features = ["temp", "atemp", "hum", "windspeed", "cnt"]
plt.figure(figsize=(6,6),dpi=70)
plt.title("Correlation of numeric features",fontsize=9)
sns.heatmap(bike_new[num_features].corr(),annot= True,cmap="mako")
plt.show()
```



13 Data Preparation for Linear Regression

```
[25]: from sklearn.preprocessing import MinMaxScaler
```

```
[26]: scaler = MinMaxScaler()
```

```
[27]: df_train.head()
```

```
[27]:
```

	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	\
483	2	1	4	0	1	1	1	18.791653	
650	4	1	10	0	0	0	1	16.126653	
212	3	0	8	0	3	1	1	31.638347	
714	4	1	12	0	1	1	2	14.862500	
8	1	0	1	0	2	1	1	5.671653	

	atemp	hum	windspeed	cnt
483	22.50605	58.7083	7.832836	6304
650	19.56980	49.4583	9.791514	7109
212	35.16460	55.0833	10.500039	4266
714	18.49690	83.8750	6.749714	3786
8	5.80875	43.4167	24.250650	822

```
[28]: df_train.columns
```

```
[28]: Index(['season', 'yr', 'mnth', 'holiday', 'weekday', 'workingday',
         'weathersit', 'temp', 'atemp', 'hum', 'windspeed', 'cnt'],
         dtype='object')
```

```
[29]: num_vars = ['temp', 'atemp', 'hum', 'windspeed', 'cnt']
df_train[num_vars] = scaler.fit_transform(df_train[num_vars])
```

```
[30]: df_train.head()
```

```
[30]:
```

	season	yr	mnth	holiday	weekday	workingday	weathersit	temp \
483	2	1	4	0	1	1	1	0.497426
650	4	1	10	0	0	0	1	0.416433
212	3	0	8	0	3	1	1	0.887856
714	4	1	12	0	1	1	2	0.378013
8	1	0	1	0	2	1	1	0.098690

	atemp	hum	windspeed	cnt
483	0.487055	0.609956	0.194850	0.722734
650	0.409971	0.513852	0.255118	0.815347
212	0.819376	0.572294	0.276919	0.488265
714	0.381804	0.871429	0.161523	0.433042
8	0.048706	0.451083	0.700017	0.092039

```
[31]: df_train.describe()
```

```
[31]:
```

	yr	holiday	workingday	temp	atemp	hum \
count	510.000000	510.000000	510.000000	510.000000	510.000000	510.000000
mean	0.501961	0.023529	0.692157	0.540901	0.515631	0.647390
std	0.500487	0.151726	0.462054	0.227898	0.213626	0.149722
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.343228	0.335807	0.536147
50%	1.000000	0.000000	1.000000	0.540519	0.525578	0.646367
75%	1.000000	0.000000	1.000000	0.740406	0.692378	0.757900
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	windspeed	cnt
count	510.000000	510.000000
mean	0.346318	0.515144

std	0.160266	0.224281
min	0.000000	0.000000
25%	0.230784	0.359468
50%	0.325635	0.516337
75%	0.434287	0.685861
max	1.000000	1.000000

```
[32]: y_train = df_train.pop('cnt')
      X_train = df_train
```

14 Use RFE to eliminate some columns

```
[33]: from sklearn.feature_selection import RFE
      from sklearn.linear_model import LinearRegression
```

```
[34]: lm = LinearRegression()
      lm.fit(X_train, y_train)
      rfe=RFE(lm,n_features_to_select=15)
      rfe = rfe.fit(X_train, y_train)
```

```
[35]: list(zip(X_train.columns,rfe.support_,rfe.ranking_))
```

```
[35]: [('season', True, 1),
      ('yr', True, 1),
      ('mnth', True, 1),
      ('holiday', True, 1),
      ('weekday', True, 1),
      ('workingday', True, 1),
      ('weathersit', True, 1),
      ('temp', True, 1),
      ('atemp', True, 1),
      ('hum', True, 1),
      ('windspeed', True, 1)]
```

```
[36]: col = X_train.columns[rfe.support_]
      col
```

```
[36]: Index(['season', 'yr', 'mnth', 'holiday', 'weekday', 'workingday',
      'weathersit', 'temp', 'atemp', 'hum', 'windspeed'],
      dtype='object')
```

```
[37]: X_train.columns[~rfe.support_]
```

```
[37]: Index([], dtype='object')
```

```
[38]: X_train_rfe = X_train[col]
```

Function to calculate VIFs and print them

```
[39]: from statsmodels.stats.outliers_influence import variance_inflation_factor
vif = pd.DataFrame()
vif['Features'] = X_train_rfe.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe.values, i) for i in
             range(X_train_rfe.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
[39]:
```

	Features	VIF
8	atemp	347.92
7	temp	335.26
9	hum	26.45
0	season	21.47
2	mnth	15.79
6	weathersit	12.91
10	windspeed	4.54
5	workingday	3.36
4	weekday	3.13
1	yr	1.98
3	holiday	1.09

```
[40]: import statsmodels.api as sm
X_train_lm1 = sm.add_constant(X_train_rfe)
lr1 = sm.OLS(y_train, X_train_lm1).fit()
```

```
C:\Users\MAHA\anaconda3\lib\site-packages\statsmodels\tsa\tsatools.py:142:
FutureWarning: In a future version of pandas all arguments of concat except for
the argument 'objs' will be keyword-only
  x = pd.concat(x[:, :order], 1)
```

```
[41]: lr1.params
```

```
[41]: const          0.271175
season          0.058342
yr              0.230654
mnth           -0.005410
holiday         -0.077059
weekday         0.003025
workingday     -0.029959
weathersit      -0.070101
temp            0.227614
atemp           0.279620
hum            -0.115018
windspeed      -0.184839
dtype: float64
```

15 Model 1 - Start with all variables selected by RFE

```
[42]: print(lr1.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  cnt    R-squared:                  0.796
Model:                        OLS    Adj. R-squared:             0.791
Method:                       Least Squares    F-statistic:              176.3
Date:                        Sat, 30 Jul 2022    Prob (F-statistic):       8.16e-164
Time:                        21:59:34    Log-Likelihood:           444.17
No. Observations:              510    AIC:                      -864.3
Df Residuals:                  498    BIC:                      -813.5
Df Model:                      11
Covariance Type:               nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.2712	0.030	8.921	0.000	0.211	0.331
season	0.0583	0.008	7.509	0.000	0.043	0.074
yr	0.2307	0.009	25.102	0.000	0.213	0.249
mnth	-0.0054	0.002	-2.205	0.028	-0.010	-0.001
holiday	-0.0771	0.031	-2.491	0.013	-0.138	-0.016
weekday	0.0030	0.002	1.327	0.185	-0.001	0.008
workingday	-0.0300	0.010	-2.931	0.004	-0.050	-0.010
weathersit	-0.0701	0.011	-6.414	0.000	-0.092	-0.049
temp	0.2276	0.142	1.602	0.110	-0.051	0.507
atemp	0.2796	0.153	1.832	0.067	-0.020	0.579
hum	-0.1150	0.042	-2.743	0.006	-0.197	-0.033
windspeed	-0.1848	0.031	-5.902	0.000	-0.246	-0.123

```

=====
Omnibus:                      59.353    Durbin-Watson:              2.002
Prob(Omnibus):                 0.000    Jarque-Bera (JB):           119.364
Skew:                         -0.671    Prob(JB):                   1.20e-26
Kurtosis:                     4.953    Cond. No.                   391.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

*Removing the variable 'atemp' based on its High p-value & High VIF

```
[43]: X_train_new = X_train_rfe.drop(["atemp"], axis = 1)
```

VIF Check


```
[44]: from statsmodels.stats.outliers_influence import variance_inflation_factor
vif = pd.DataFrame()
vif['Features'] = X_train_new.columns
vif['VIF'] = [variance_inflation_factor(X_train_new.values, i) for i in
             range(X_train_new.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
[44]:
```

	Features	VIF
8	hum	25.94
0	season	21.40
2	mnth	15.78
6	weathersit	12.89
7	temp	7.87
9	windspeed	4.44
5	workingday	3.34
4	weekday	3.13
1	yr	1.98
3	holiday	1.09

```
[45]: X_train_lm2 = sm.add_constant(X_train_new)
lr2 = sm.OLS(y_train, X_train_lm2).fit()
```

C:\Users\MAHA\anaconda3\lib\site-packages\statsmodels\tsa\tsatools.py:142:
FutureWarning: In a future version of pandas all arguments of concat except for
the argument 'objs' will be keyword-only
x = pd.concat(x[:, :order], 1)

```
[46]: lr2.params
```

```
[46]: const          0.277236
season           0.058907
yr              0.230812
mnth            -0.005525
holiday         -0.077668
weekday          0.002778
workingday      -0.028730
weathersit       -0.070773
temp            0.484720
hum            -0.110721
windspeed       -0.195338
dtype: float64
```

16 Model-2 bulid after removing atemp:

```
[47]: print(lr2.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          cnt      R-squared:                0.794
Model:                  OLS      Adj. R-squared:           0.790
Method:                 Least Squares      F-statistic:        192.7
Date:                   Sat, 30 Jul 2022      Prob (F-statistic):    3.03e-164
Time:                   21:59:35      Log-Likelihood:        442.45
No. Observations:       510      AIC:                   -862.9
Df Residuals:           499      BIC:                   -816.3
Df Model:                10
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.2772	0.030	9.153	0.000	0.218	0.337
season	0.0589	0.008	7.570	0.000	0.044	0.074
yr	0.2308	0.009	25.061	0.000	0.213	0.249
mnth	-0.0055	0.002	-2.248	0.025	-0.010	-0.001
holiday	-0.0777	0.031	-2.505	0.013	-0.139	-0.017
weekday	0.0028	0.002	1.218	0.224	-0.002	0.007
workingday	-0.0287	0.010	-2.810	0.005	-0.049	-0.009
weathersit	-0.0708	0.011	-6.464	0.000	-0.092	-0.049
temp	0.4847	0.022	21.899	0.000	0.441	0.528
hum	-0.1107	0.042	-2.639	0.009	-0.193	-0.028
windspeed	-0.1953	0.031	-6.330	0.000	-0.256	-0.135

```
=====
Omnibus:                 52.728      Durbin-Watson:           2.002
Prob(Omnibus):            0.000      Jarque-Bera (JB):        103.061
Skew:                     -0.613      Prob(JB):                4.17e-23
Kurtosis:                 4.829      Cond. No.                 92.6
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1) Removing the variable 'hum' based on its Very High 'VIF' value. 2) Even though the VIF of hum is second highest, we decided to drop 'hum' and not 'temp' based on general knowledge that temperature can be an important factor for a business like bike rentals, and wanted to retain 'temp'.

```
[48]: X_train_new = X_train_new.drop(["hum"], axis = 1)
```

VIF check

```
[49]: from statsmodels.stats.outliers_influence import variance_inflation_factor
vif = pd.DataFrame()
vif['Features'] = X_train_new.columns
vif['VIF'] = [variance_inflation_factor(X_train_new.values, i) for i in
             range(X_train_new.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
[49]:
```

	Features	VIF
0	season	21.36
2	mnth	15.37
7	temp	6.75
6	weathersit	5.95
8	windspeed	4.43
5	workingday	3.09
4	weekday	3.08
1	yr	1.98
3	holiday	1.08

```
[50]: X_train_lm3 = sm.add_constant(X_train_new)
lr3 = sm.OLS(y_train, X_train_lm3).fit()
```

C:\Users\MAHA\anaconda3\lib\site-packages\statsmodels\tsa\tsatools.py:142:
FutureWarning: In a future version of pandas all arguments of concat except for
the argument 'objs' will be keyword-only
x = pd.concat(x[:, :order], 1)

```
[51]: lr3.params
```

```
[51]: const          0.232782
season          0.059786
yr             0.233247
mnth          -0.006415
holiday        -0.079728
weekday         0.002890
workingday     -0.031422
weathersit     -0.088825
temp           0.472745
windspeed     -0.171276
dtype: float64
```

17 Model-3 bulid after removing hum:

```
[52]: print(lr3.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          cnt      R-squared:                0.791
Model:                  OLS      Adj. R-squared:           0.788
Method:                 Least Squares      F-statistic:         210.8
Date:                   Sat, 30 Jul 2022    Prob (F-statistic):    6.34e-164
Time:                   21:59:36           Log-Likelihood:        438.92
No. Observations:       510             AIC:                  -857.8
Df Residuals:           500             BIC:                  -815.5
Df Model:                9
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.2328	0.025	9.193	0.000	0.183	0.283
season	0.0598	0.008	7.644	0.000	0.044	0.075
yr	0.2332	0.009	25.303	0.000	0.215	0.251
mnth	-0.0064	0.002	-2.619	0.009	-0.011	-0.002
holiday	-0.0797	0.031	-2.557	0.011	-0.141	-0.018
weekday	0.0029	0.002	1.260	0.208	-0.002	0.007
workingday	-0.0314	0.010	-3.070	0.002	-0.052	-0.011
weathersit	-0.0888	0.009	-10.329	0.000	-0.106	-0.072
temp	0.4727	0.022	21.692	0.000	0.430	0.516
windspeed	-0.1713	0.030	-5.775	0.000	-0.230	-0.113

```

=====
Omnibus:                50.579      Durbin-Watson:           1.991
Prob(Omnibus):           0.000      Jarque-Bera (JB):        98.220
Skew:                    -0.593      Prob(JB):                4.70e-22
Kurtosis:                4.793      Cond. No.                 65.1
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

*Removing the variable 'season' based on its Very High 'VIF' value

```
[53]: X_train_new = X_train_new.drop(["season"], axis = 1)
```

VIF check

```
[54]: from statsmodels.stats.outliers_influence import variance_inflation_factor
vif = pd.DataFrame()
vif['Features'] = X_train_new.columns
vif['VIF'] = [variance_inflation_factor(X_train_new.values, i) for i in
↳ range(X_train_new.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
```

```
vif
```

```
[54]:      Features    VIF
      6      temp  5.95
      5  weathersit  5.91
      1      mnth  4.51
      7  windspeed  4.43
      3    weekday  3.08
      4  workingday  3.08
      0         yr  1.97
      2    holiday  1.07
```

```
[55]: X_train_lm4 = sm.add_constant(X_train_new)
      lr4 = sm.OLS(y_train, X_train_lm4).fit()
```

```
C:\Users\MAHA\anaconda3\lib\site-packages\statsmodels\tsa\tsatools.py:142:
FutureWarning: In a future version of pandas all arguments of concat except for
the argument 'objs' will be keyword-only
      x = pd.concat(x[:, :order], 1)
```

```
[56]: lr4.params
```

```
[56]: const      0.259831
      yr         0.235243
      mnth       0.008893
      holiday    -0.065440
      weekday     0.002748
      workingday -0.030299
      weathersit  -0.088357
      temp       0.517163
      windspeed  -0.182014
      dtype: float64
```

18 Model-4 build after removing Season:

```
[57]: print(lr4.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          cnt      R-squared:                0.767
Model:                  OLS      Adj. R-squared:           0.763
Method:                 Least Squares      F-statistic:        206.2
Date:                  Sat, 30 Jul 2022      Prob (F-statistic):    3.84e-153
Time:                  21:59:36      Log-Likelihood:        410.73
No. Observations:      510      AIC:                  -803.5
Df Residuals:          501      BIC:                  -765.4
```

```

Df Model:                8
Covariance Type:         nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const          0.2598        0.026        9.816      0.000        0.208        0.312
yr             0.2352        0.010       24.182      0.000        0.216        0.254
mnth           0.0089        0.001        5.973      0.000        0.006        0.012
holiday        -0.0654        0.033       -1.992      0.047       -0.130       -0.001
weekday         0.0027        0.002        1.135      0.257       -0.002        0.008
workingday     -0.0303        0.011       -2.804      0.005       -0.052       -0.009
weathersit      -0.0884        0.009       -9.732      0.000       -0.106       -0.071
temp           0.5172        0.022       23.321      0.000        0.474        0.561
windspeed      -0.1820        0.031       -5.819      0.000       -0.243       -0.121
=====
Omnibus:                44.143   Durbin-Watson:                1.983
Prob(Omnibus):           0.000   Jarque-Bera (JB):        68.248
Skew:                   -0.604   Prob(JB):                1.51e-15
Kurtosis:                4.324   Cond. No.                61.8
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

19 Final Model Interpretation

Hypothesis Testing:

Hypothesis testing states that:

$H_0: B_1 = B_2 = \dots = B_n = 0$

H_1 : at least one $B_i \neq 0$

lr4 model coefficient values const 0.259831 yr 0.235243 mnth 0.008893 holiday -0.065440 weekday 0.002748 workingday -0.030299 weathersit -0.088357 temp 0.517163 windspeed -0.182014 dtype: float64

20 Interpret

From the lr6 model summary, it is evident that all our coefficients are not equal to zero which means We REJECT the NULL HYPOTHESIS

The equation of best fitted surface based on model lr4:

$$\text{cnt} = 0.2598 + (\text{yr} \times 0.2352) + (\text{mnth} \times 0.0089) - (\text{workingday} \times 0.0303) - (\text{weathersit} \times 0.0884) + (\text{temp} \times 0.5172) - (\text{windspeed} \times 0.1820)$$

Interpretation of Coefficients: yr: A coefficient value of '0.2352' indicated that a unit increase in yr variable, increases the bike hire numbers by 0.2352 units.

mnth: A coefficient value of '0.0089' indicated that a unit increase in mnth variable, increase the bike hire numbers by 0.0089 units.

yr: A coefficient value of '0.2308' indicated that a unit increase in yr variable, increases the bike hire numbers by 0.2308 units.

workingday: A coefficient value of '-0.0303' indicated that a unit increase in workingday variable decrease the bike hire numbers by 0.0303 units.

weathersit: A coefficient value of '-0.0884' indicated that, a unit increase in windspeed variable decreases the bike hire numbers by 0.0884 units.

temp: A coefficient value of '0.5172' indicated that, a unit increase in workingday variable increases the bike hire numbers by 0.5172 units.

windspeed: A coefficient value of '-0.1820' indicated that, a unit increase in workingday variable decrease the bike hire numbers by 0.1820 units.

21 Error terms are normally distributed with mean zero (not X, Y)

Residual Analysis Of Training Data


```
[58]: y_train_pred = lr4.predict(X_train_lm4)
```

```
[59]: res = y_train-y_train_pred
      #plot histogram error term
      fig = plt.figure()
      sns.distplot((res), bins = 20)
      fig.suptitle('Error Terms', fontsize = 20)
      plt.xlabel('Errors', fontsize = 18)
```

C:\Users\MAHA\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
[59]: Text(0.5, 0, 'Errors')
```



22 Interpret:

From the above histogram, we could see that the Residuals are normally distributed. Hence our assumption for Linear Regression is valid.

23 MAKING PREDICTION USING FINAL MODEL

Now that we have fitted the model and checked the assumptions, it's time to go ahead and make predictions using the final model (lr4)

Applying the scaling on the test sets

```
[60]: num_vars = ['temp', 'atemp', 'hum', 'windspeed', 'cnt']
      df_test[num_vars] = scaler.transform(df_test[num_vars])
```

```
[61]: df_test.head()
```

```
[61]:
```

	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	\
22	1	0	1	0	2	1	1	0.046591	
468	2	1	4	0	0	0	1	0.543115	
553	3	1	7	0	1	1	1	0.951196	
504	2	1	5	0	1	1	1	0.699909	


```
353      4    0   12      0      4      1      2  0.407087
```

```
      atemp      hum  windspeed      cnt
22  0.025950  0.453529  0.462217  0.110907
468 0.536771  0.522511  0.347424  0.855729
553 0.933712  0.596104  0.212829  0.534975
504 0.662746  0.551083  0.478229  0.817648
353 0.416610  0.618615  0.080770  0.428900
```

```
[62]: df_test.describe()
```

```
[62]:
```

	yr	holiday	workingday	temp	atemp	hum \
count	219.000000	219.000000	219.000000	219.000000	219.000000	219.000000
mean	0.493151	0.041096	0.689498	0.551225	0.527528	0.662567
std	0.501098	0.198967	0.463759	0.229463	0.215434	0.143562
min	0.000000	0.000000	0.000000	0.046591	0.025950	0.301299
25%	0.000000	0.000000	0.000000	0.356479	0.348019	0.553031
50%	0.000000	0.000000	1.000000	0.557653	0.549198	0.662338
75%	1.000000	0.000000	1.000000	0.751309	0.709163	0.762338
max	1.000000	1.000000	1.000000	0.984424	0.980934	1.010390

	windspeed	cnt
count	219.000000	219.000000
mean	0.346706	0.518889
std	0.159553	0.219953
min	0.073090	0.055683
25%	0.232689	0.364703
50%	0.328208	0.525771
75%	0.435708	0.676887
max	0.824380	0.963300

Dividing into X_test and y_test

```
[63]: y_test = df_test.pop('cnt')
      X_test = df_test
      X_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 219 entries, 22 to 313
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   season      219 non-null    category
1   yr          219 non-null    int64
2   mnth       219 non-null    category
3   holiday     219 non-null    int64
4   weekday     219 non-null    category
5   workingday  219 non-null    int64
```

```

6  weathersit    219 non-null    category
7  temp         219 non-null    float64
8  atemp        219 non-null    float64
9  hum          219 non-null    float64
10 windspeed    219 non-null    float64
dtypes: category(4), float64(4), int64(3)
memory usage: 15.6 KB

```

```

[64]: col1=X_train_new.columns
      X_test=X_test[col1]
      X_test_lm4 = sm.add_constant(X_test)
      X_test_lm4.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 219 entries, 22 to 313
Data columns (total 9 columns):
 #   Column        Non-Null Count  Dtype
---  -
0   const         219 non-null    float64
1   yr            219 non-null    int64
2   mnth         219 non-null    category
3   holiday       219 non-null    int64
4   weekday       219 non-null    category
5   workingday    219 non-null    int64
6   weathersit     219 non-null    category
7   temp          219 non-null    float64
8   windspeed     219 non-null    float64
dtypes: category(3), float64(3), int64(3)
memory usage: 13.5 KB

```

```

C:\Users\MAHA\anaconda3\lib\site-packages\statsmodels\tsa\tsatools.py:142:
FutureWarning: In a future version of pandas all arguments of concat except for
the argument 'objs' will be keyword-only
    x = pd.concat(x[:, :order], 1)

```

```

[65]: y_pred = lr4.predict(X_test_lm4)

```

24 MODEL EVALUATION

```


[66]: fig = plt.figure()
      plt.scatter(y_test, y_pred, alpha=.5)
      fig.suptitle('y_test vs y_pred', fontsize = 20)
      plt.xlabel('y_test', fontsize = 18)
      plt.ylabel('y_pred', fontsize = 16)
      plt.show()

```

```

# Plot heading
# X-label

```



25 R^2 Value for TEST

```
[67]: from sklearn.metrics import r2_score  
r2_score(y_test, y_pred)
```

```
[67]: 0.7658194327343915
```

26 Adjusted R^2 Value for TEST

```
[68]: r2=0.7658194327343915
```

```
[69]: X_test.shape
```

```
[69]: (219, 8)
```

```
[70]: n = X_test.shape[0]  
p = X_test.shape[1]  
adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)  
adjusted_r2
```

[70]: 0.7568982682671301

27 Final Result Comparison

Train R^2 :0.767

Train Adjusted R^2 :0.763

Test R^2 :0.766

Test Adjusted R^2 :0.757

This seems to be a really good model that can very well ‘Generalize’ various datasets.

28 FINAL REPORT

As per our final Model, the top 3 predictor variables that influences the bike booking are:

temp: A coefficient value of ‘0.5172’ indicated that, a unit increase in workingday variable increases the bike hire numbers by 0.5172 units.

yr: A coefficient value of ‘0.2308’ indicated that a unit increase in yr variable, increases the bike hire numbers by 0.2308 units.

windspeed: A coefficient value of ‘-0.1820’ indicated that, a unit increase in workingday variable decrease the bike hire numbers by 0.1820 units.

[]: